

The Copenhagen cholera epidemics 1853 in contemporary Danish newspaper prose: Some back-of-an-envelope calculations made during April 2020

Sigfrid Lundberg

The Royal Danish Library

Introduction

Inspired by the fact that we are all hiding away from the Covid-19 I wanted to take a closer look at some other epidemic. My hope was to find patterns of change in language use mirroring sentiments and attitudes expressed in words, bigrams and trigrams and frequency distributions. My first impulse was to analyze the Spanish flu 1918, but since we have as yet little data in our public corpora from that period I turned my attention to the Cholera epidemic in Copenhagen 1853. At the time, Copenhagen had around 130.000¹ inhabitants out of which 7.219 were diagnosed as having the disease out of which 4.737 (56,7%) died. Christian Molbech writes in a letter to Christian Knud Frederik Molbech, the July 3, 1853:

The first cholera case was found as early as the 12th June (i.e., the day before the opening of the parliament in June) — a young man from the area close to the naval barracks, who for quite some time had worked on a dredge. He has recovered, and was released from hospital the 25th. The authorities were able to conceal the earliest few cases for a fortnight.²




Figure 1. Number of newspaper pages mentioning any of the words *kolera*, *cholera*, *epidemi* or *pest*.

¹ <http://www2.kb.dk/udstillinger/medhist/kolera/koebenhavn1853.htm>

² Translated by the author. <https://bit.ly/3cOzIAg>

The epidemic spread from Copenhagen to other areas. Outside the capital 1.951 fatal cases were reported.³

The spread of an infectious disease does not occur in a vacuum; the Copenhagen epidemic is a part of the cholera pandemic (1846-1860)⁴ which is the third out of seven global outbreaks.⁵

Here I present some data analyses: (1) The number of newspaper articles mentioning *epidemi*, *pest*, *kolera* and *cholera* through the period 1846 to 1860 (Figure 1). (2) I repeated that analysis with the frequencies of words rather than articles in weekly aggregations of texts. (3) Finally I visualize a set of bigrams extracted from the corpus.

The data

The data was retrieved from Royal Danish Library's LOAR Repository.⁶ It comes in csv format, aggregated to yearly chunks. Each line is basically corresponding to one page of text in the newspaper.

Some low-hanging fruits

Using the Unix `grep` command makes it easy to count the number of pages containing a given word. So the graph in Figure 1, is simple plot of what is extracted using the shell script in Figure 2.

```
#!/bin/sh
echo "#year kolera cholera epidemi pest farsot"
for file in artikler*
do
    year=$(echo $file | tr -d '[:alpha:][:punct:]')
    kolera=$(grep -i kolera $file | wc -l)
    cholera=$(grep -i cholera $file | wc -l)
    epidemi=$(grep -i epidemi $file | wc -l)
    pest=$(grep -i pest $file | wc -l)
    farsot=$(grep -i farsot $file | wc -l)
    echo "$year $kolera $cholera $epidemi $pest $farsot"
done
```

Figure 2. Script for extracting the data for plotting Figure 1

Three things spring to my mind: (1) The obvious one. People wrote about this cholera outbreak in Copenhagen 1853. They wrote a lot. (2) The effect on the public discourse of the epidemic lasted for years and it did not return to pre-epidemic levels that decade. Almost certainly the people felt like we do: Will it ever be the same again? That is based on the fact that the article count in Figure 1 never really decreases to counts before the cholera outbreak. (3) People hadn't settled on how to spell the name of the disease yet, cholera and kolera. (4) The word *pest* seems to be more widely used in Danish at the time, possibly as something annoying but also as epidemic disease in general.

³ https://da.wikipedia.org/wiki/Koleraepidemien_i_K%C3%B8benhavn_1853

⁴ https://en.wikipedia.org/wiki/1846%E2%80%931860_cholera_pandemic

⁵ https://en.wikipedia.org/wiki/Cholera_outbreaks_and_pandemics

⁶ Newspapers from Royal Danish Library <https://loar.kb.dk/handle/1902/157>

I also made another analysis, where I aggregated text by week, and used text tokenized by word, see Figure 3. The week numbering starts with 1 at 1846-01-01 and ends at 1860-12-31 with week 781. The huge peak in the graph starts at 1853-06-24 at week 389. It reaches its highest point at the end of July, 1853-07-29, week 395. From what I can tell it continues until the beginning of October, but the mentions of *cholera* in the corpus is higher than the pre-outbreak into March the following year, week 425. There is a discrepancy between the two analyses, in that the article level analyses (Figure 1) implies a longer effect of the outbreak on the discourse than the per word one (Figure 3.)




Figure 3. Here I use tokenized text, i.e., there is only one word per line, and I also aggregated the texts by week.

Bigrams and trigrams

N-grams generation is not an end in itself. They are a starting point for many kinds of statistical and machine learning methods and applications in natural language processing.⁷ You extract n-grams by "sliding" through your tokenized text word by word, and sample n words. Bigrams and trigrams are special cases where you sample two or three words. I am a great fan of simplicity and elegance and hence I love Kenneth Ward Church's tool box which he described in his wonderful classic *Unix for Poets*⁸ where you are given methods to do a text analyses using simple tools that you find on all Unix/Linux machines. Here are some examples of trigrams containing the word *hovedstaden* across the weeks, i.e., the capital (in decreasing frequency):

```
cholera i hovedstaden
hovedstaden bortriver cholera
hovedstaden er cholera
hovedstaden af cholera
hovedstaden forekomne cholera
choleraepidemien i hovedstaden
epidemien i hovedstaden
```

⁷ <https://en.wikipedia.org/wiki/N-gram>

⁸ <http://doc.cat-v.org/unix/for-poets/>

hovedstaden udbrudte cholera

For some reason, I don't understand why, the actual name of the city Copenhagen, København (or whatever spelling) seem to occur rarely in comparison with the word hovedstaden, which is frequent, in particular 1853. Here is another sequence (containing *cholera er*) and sampled such that they begin with that phrase:

cholera er af
cholera er afgaaet
cholera er aften
cholera er aldeles
cholera er alter
cholera er aner
cholera er at
cholera er atter
cholera er begjeert
cholera erboldes som
cholera er borlfalden
cholera er bortfalden
cholera er cadetfiibs
cholera ere altsaa
cholera ere berovede

The rest of this paper consists of a visualizations

1. of the third word in trigrams starting with *kolera* or *cholera i*. The third word is more often than not a place name. Repetitions are excluded, so the heights of the bars are more measures of diversity rather than abundance. See Figure 4.
2. of bigrams containing any of the words *epidemi*, *kolera* or *cholera*.

The bigrams are shown as bars in a set of bar diagrams, one diagram per eight weeks interval. It comprises about a year, from 1853-05-27 (my week 386) to 1854-05-06 (my week 435). When the epidemics is worst, the number of bigrams is so high that I decided to **take only the 50 most frequent ones**, or the text would not be readable. Then all the bars are of equal length, obviously. The bars look somewhat like this, the number is the absolute frequency of the following bigram that week.

46 af cholera
42 kolera den
36 kolera ten
36 a kolera
23 kolera hen
15 kolera
10 kolera drn
9 sidste epidemi
9 kolera i
9 kolera d




Figure 4. Trigrams per year with the text beginning with cholera i. The highest bar is slightly massaged. I removed repetitions and incomprehensible OCR errors in order to make it fit into the page.

Lessons learned

- 1 My idea was that just inspection of the words in their context should give information about emotions or sentiments. I suppose that this was just due to my naivete when it comes to languages, since I cannot really find anything of that kind.
- 2 The tools coming with a standard Linux distribution isn't as unicode compliant one would wish. Some twenty years later, the nice utility `tr` is still not really UTF-8. Neither is the GNU implementation of Brian Kernighans graphics language `pic` which I used for the bigrams.
- 3 There is another big problem connected with the use of our newspaper text corpora for computational linguistics. The poor OCR quality.